

Latency

Andrew Kondratovich
andrew.kondratovich@gmail.com

*«**Latency** is a time interval between the stimulation and response, or, from a more general point of view, is a time delay between the cause and the effect of some change in the system being observed»*

— <https://en.wikipedia.org/>

Latency versus bandwidth

- Latency is not the same as bandwidth.
- Reducing latency does not always improve bandwidth, in many cases it has the reverse effect.
- It still takes the same amount of time to get the first bit but if you need 10kb before you can start processing then bandwidth matters.

Infinite bandwidth, zero latency


$$C = 299,792,458 \text{ m/s}$$

- Every 300m of cable/fibre/microwave adds at least $1\mu\text{s}$ of latency.
- 300km will add at least 1ms.
- 1ms is worth about \$100m in the trading world.


Latency behaviour

- Each operation occurrence has its own latency.
- What we care about is how latency *behaves*.
- Behavior is more than «what was the common case?»


Response time over load



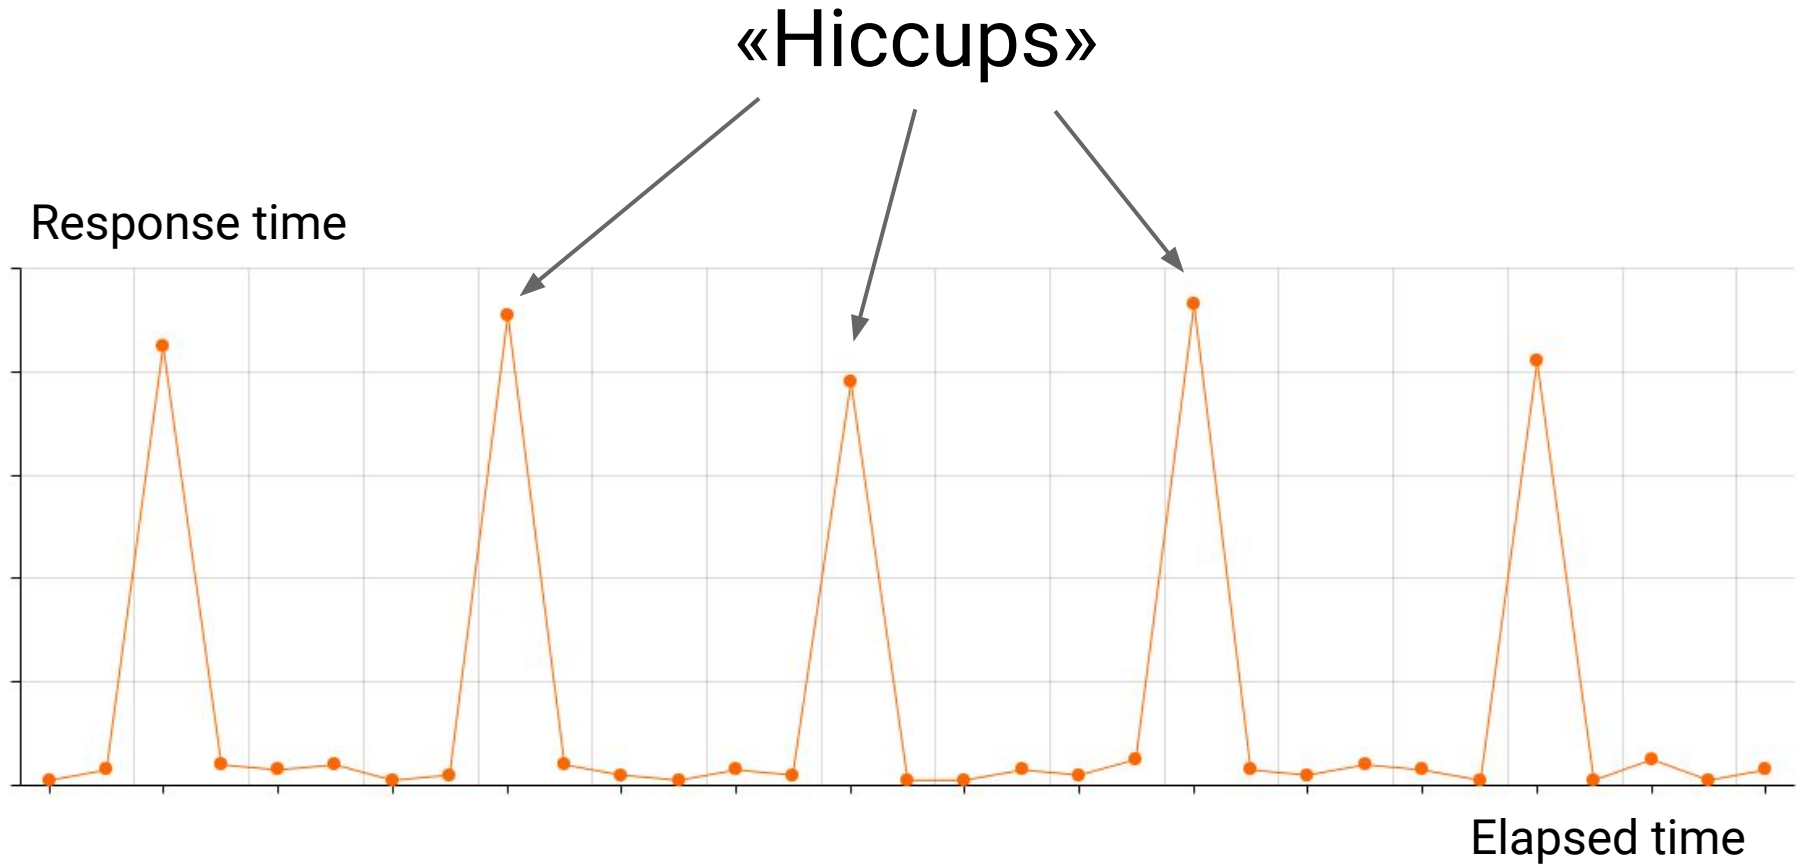
Response time over load



Response time over time



Response time over time



Hiccups

- They don't look like a normal distribution.
- They usually look like periodic freezes.
- A complete shift from one state to another:
 1. everything is okay
 2. something gone bad
 3. disaster

Deal with it

- What do we want the latency to be?
- Different applications have different needs.
- Requirements should reflect application needs.
- Better way to deal with hiccups is measuring ***percentiles***.

What do we care about?

- Care about the worst case?
- Care about the 99.99%?
- Care if 1% of operations fail?
- Care about few fastest events?

...

Latency Numbers

L1 cache reference	0.5 ns		
Branch mispredict	5 ns		
L2 cache reference	7 ns		
Mutex lock/unlock	25 ns		
Main memory reference	100 ns		
Compress 1K bytes with Zippy	3,000 ns	=	3 µs
Send 2K bytes over 1 Gbps network	20,000 ns	=	20 µs
SSD random read	150,000 ns	=	150 µs
Read 1 MB sequentially from memory	250,000 ns	=	250 µs
Round trip within same datacenter	500,000 ns	=	0.5 ms
Read 1 MB sequentially from SSD	1,000,000 ns	=	1 ms
Disk seek	10,000,000 ns	=	10 ms
Read 1 MB sequentially from disk	20,000,000 ns	=	20 ms
Send packet CA->Netherlands->CA	150,000,000 ns	=	150 ms

Java memory 1/4

- Java is very inefficient at storing data in memory.
- Objects get created everywhere, they're good for OOP but crap if you're trying to get performance out of your machine.
- Each object created will need to be garbage collected.

Java memory 2/4

- Objects have 12/16 bytes headers.
- Java bloating is endemic: *String*, *Double*, *BigDecimal*, *Date*, *LinkedList*, *Iterator*.

```
public class Quote {  
    private Date tradeDate;  
    private BigDecimal bid;  
    private String currency;  
    ...  
}
```

- Abstraction costs dearly.

Java memory 3/4

- We can use **compression**, but it's slow.
- We want **compaction** not compression.

```
public class Quote {  
    private byte[] data;  
}
```

- Just one object, fast to allocate.
- If we can encode the data in the binary then it's fast to query too.

Java memory 4/4

- Identical API.
- We can use ByteBuffer instead of byte array.

```
public Date getTradeDate() {  
    return tradeDate;  
}
```

```
public Date getTradeDate() {  
    return new Date((((long) data[0] & 0xff) << 0) |  
                    (((long) data[1] & 0xff) << 8) |  
                    ...);  
}
```

Optimizations 1/3

- Preallocation

Queues, buffers, networking, etc...

- Redundant write elimination

```
elements[index] = item;  
// versus  
if (elements[index] != item) {  
    elements[index] = item;  
}
```

Optimizations 2/3

- Passive and Active waiting

Parking versus Spinning

- Backoff

Thread.yield();

- Thread affinity

Thread context relocation workaround

- Sequential and random access

ArrayList and LinkedList

- True and false sharing

Optimizations 3/3

– CAS/TAS and TTAS

– Atomic.lazySet()

Non-volatile write into volatile variable

volatile store

lazySet

`mbar(store|store)`

`store a`

`mbar(store|load)`

`mbar(store|store)`

`store a`

Q/A